

OpenClaw

Done-For-You AI Agent Installation

Your private AI workforce, running on dedicated hardware.
No cloud dependency. No per-seat fees. Unlimited usage.
Starting at \$5,000.

Instant Access

Start using OpenClaw the same day you place your deposit

Dedicated Hardware

Your own Apple Mac Mini ships pre-configured and ready to run

Full Ownership

The machine, the models, and the data are permanently yours

Prepared by Organized AI
clawdbot.organizedai.vip

How It Works

OpenClaw is a done-for-you AI agent that runs entirely on dedicated hardware. Unlike cloud-based AI tools that charge per seat, per message, or per token, OpenClaw gives you unlimited AI usage on a machine you own outright. No recurring API fees. No data leaving your network. Every skill and plugin comes from the Organized AI Marketplace — vetted, tested, and ready for production.

The Test-Drive-to-Ownership Model

- 1 Place Your Deposit**
You secure your slot with a deposit. The same day, we provision your OpenClaw agent on our high-performance infrastructure and hand you the keys.
- 2 Start Using OpenClaw Immediately**
Your agent is live on our Mac Studio M3 Ultra with 256 GB. Chat via WhatsApp, web, or API. Run tasks, generate content, analyze data — full power, zero per-use cost.
- 3 We Configure Your Dedicated Mac Mini**
While you use the test drive, we order, configure, and stress-test your own Apple Mac Mini with your specific plugins, models, and preferences from the Organized AI Marketplace.
- 4 Your Mac Mini Ships to You**
Once your dedicated hardware is ready and verified, we ship it to your location. We migrate your data, preferences, and conversation history over seamlessly.
- 5 Ongoing Support + Updates**
We provide remote maintenance, model updates, plugin additions from the Organized AI Marketplace, and priority support. Your machine gets better over time.

Why This Model?

Most AI services ask you to wait for setup before you see any value. We flip that. Your deposit gets you working with OpenClaw on day one — running on enterprise-grade hardware with the full model stack. The test drive is not a demo: it is the real product, running real models, doing real work. By the time your dedicated Mac Mini arrives, you already know exactly what you are getting because you have been using it.

What You Get

Test Drive: Enterprise-Grade Inference

During your test-drive period, your agent runs on our Mac Studio M3 Ultra — the most powerful desktop Apple has ever built, with 256 GB of unified memory and 819 GB/s bandwidth. This gives you access to the full range of AI models including large 70B-parameter reasoning models that deliver deep analysis and complex planning. You experience the best of what local AI can offer from day one.

Your Hardware: Apple Mac Mini (or Mac Studio Upgrade)

Your dedicated OpenClaw installation ships on an Apple Mac Mini — a compact, silent, energy-efficient machine configured specifically for your workloads. It runs the AI models you use day-to-day, with intelligent routing that selects the right model for each task automatically. Need more power? Upgrade to a Mac Studio M3 Ultra with 256 GB of unified memory for the full 70B-parameter model stack running locally.

Specification	Mac Mini (Standard)	Mac Studio M3 Ultra (Upgrade)
Processor	Apple M4 Pro	Apple M3 Ultra — 32-core CPU
Unified Memory	Up to 64 GB	256 GB — runs 70B models natively
Best For	3B–32B models, daily workflows	Full model stack including 70B reasoning
Noise / Power	Near silent, ~40W (~\$5/mo)	Near silent, ~150W (~\$13/mo)
Form Factor	5 x 5 x 2 in — palm-sized	7.7 x 7.7 x 3.7 in — compact desktop

The AI Models: Running Locally via ClawRouter

Your Mac Mini runs multiple AI models simultaneously. An intelligent routing layer called ClawRouter automatically selects the right model for each task — a fast model for quick questions, a specialized model for code, and a capable reasoning model for analysis. You never have to think about which model to use. The system decides in under a millisecond.

Task Type	Model	Response Speed	Best For
Quick Tasks	3B parameter	Instant	Q&A, lookups, greetings, triage
Standard Work	8B parameter	Fast	Writing, analysis, code review

Deep Reasoning	14B–32B parameter	Thoughtful	Strategy, planning, complex reports
Code Generation	7B specialized	Fast	Writing and debugging code

The Organized AI Marketplace

OpenClaw is not a blank chatbot. It ships with pre-built skills and plugins from the **Organized AI Marketplace** — a vetted library of production-ready capabilities developed and tested for real business use. Every plugin in the marketplace has been reviewed for quality, security, and reliability. We select the right modules during your test drive based on what you actually need, and you can add more at any time.

Capabilities by Vertical

Marketing

Brand voice enforcement, campaign planning, content creation across all channels, competitive analysis, performance analytics, and SEO auditing.

Sales

Account research, call preparation with attendee briefings, competitive battlecards, personalized outreach drafting, pipeline reviews, and daily sales briefings.

Product Management

Feature spec writing, roadmap prioritization using RICE/MoSCoW/ICE frameworks, metrics tracking with OKR support, stakeholder communications, and user research synthesis.

Data & Analytics

Dataset exploration and profiling, statistical analysis, SQL query generation for any warehouse dialect, interactive dashboard creation, and data validation before sharing.

Developer Tools

Codebase organization, project scaffolding, phased implementation planning, Stripe integration best practices, and frontend design generation.

Advertising & MarTech

Google Tag Manager automation, LinkedIn CAPI server-side tracking, Meta Ads account auditing, and cross-platform conversion infrastructure.

**70+ vetted plugins across 6 verticals. Every capability runs locally on your hardware.
Your data, prompts, and outputs never leave your machine.**

The Economics

Cloud AI vs. OpenClaw: A Cost Comparison

Cloud AI services charge per user, per message, or per API call. Costs scale linearly with usage and team size. OpenClaw costs are fixed — the Mac Mini runs whether you send one message or ten thousand.

Cost Category	Cloud AI (typical)	OpenClaw
5-seat team, monthly	\$500 – \$2,000/mo	\$0 after purchase
API usage (heavy)	\$1,000 – \$5,000/mo	\$0 — runs locally
Data privacy compliance	Complex / vendor-dependent	Built in — data stays local
Usage limits	Rate limits, token caps	Unlimited
Year 1 total cost (est.)	\$18,000 – \$84,000	Hardware + support plan
Year 2+ cost	Same as Year 1	Support plan only

What the Hardware Investment Covers

The Apple Mac Mini is a one-time purchase that becomes a permanent asset on your balance sheet. Unlike SaaS subscriptions that disappear when you stop paying, this hardware and everything on it belongs to you.

- Apple Mac Mini (standard) or Mac Studio M3 Ultra (upgrade) — pre-configured and stress-tested
- Pre-installed OpenClaw agent platform with ClawRouter intelligent model routing
- AI models downloaded, benchmarked, and optimized for your specific workloads
- Curated plugins and skills from the Organized AI Marketplace tailored to your vertical
- Sandboxed execution environment for safe code and command running
- Secure remote access via encrypted mesh networking (Tailscale)
- Full data migration from your test-drive period on our M3 Ultra
- Ongoing remote support, model updates, and new marketplace plugin additions

Your Journey

Phase	What Happens	What You Experience
Deposit	We provision your agent on our Mac Studio M3 Ultra and configure your initial plugins from the Organized AI Marketplace	Same-day access to your live OpenClaw agent via WhatsApp, web, or API — running on enterprise-grade hardware
Test Drive	You use OpenClaw for real work on our infrastructure. We observe usage patterns, refine model routing, and add marketplace plugins you need	A working AI agent that gets smarter and more useful as we dial in your preferences and vertical fit
Hardware Config	We order your Mac Mini, install macOS and OpenClaw, download and benchmark models, install your marketplace plugins, and configure networking	No disruption — you keep using your test-drive agent on our M3 Ultra the entire time
Handoff	We ship the Mac Mini, migrate your data and conversation history, verify the installation, and hand over documentation	Seamless transition. Same agent, same capabilities, now running on your own dedicated hardware
Ongoing Support	Remote monitoring, model updates, new marketplace plugin additions, troubleshooting, and priority support	Your agent keeps improving. New models, new plugins from the marketplace, zero effort on your end

Security and Privacy

OpenClaw is built with a defense-in-depth security model. Every component operates under deny-by-default permissions. All AI inference happens locally — your prompts, business data, and generated outputs never touch a third-party server. The machine communicates over an encrypted mesh network for remote access and support, with no open ports exposed to the internet.

- All AI inference runs on-device — no data sent to OpenAI, Anthropic, or any cloud provider
- Encrypted mesh networking (Tailscale) for remote access — no open ports
- Sandboxed execution environment — agent commands run in isolation
- Per-client data isolation — your data is separated at the filesystem level
- Deny-by-default permissions — the agent must be explicitly granted access to every resource

Service Plans

OpenClaw is a one-time installation, not a subscription. The service starts at **\$5,000** and scales up to **\$35,000** depending on how many marketplace skills, local models, and hardware power you need. Every tier includes the test-drive period on our M3 Ultra infrastructure, a dedicated machine, installation, and ongoing support.

What Drives the Price

Component	Starting at \$5,000	Up to \$35,000
Hardware	Apple Mac Mini (M4 Pro)	Apple Mac Studio M3 Ultra (256 GB)
Test-drive on M3 Ultra	✓	✓
ClawRouter model routing	✓	✓
Local AI models	Core models (3B–14B)	Full stack including 70B reasoning
Organized AI Marketplace skills	Select vertical	All 70+ vetted plugins across 6 verticals
WhatsApp / web / API access	✓	✓
MCP tool discovery	–	✓
Sandbox execution	–	✓
Observability dashboard	–	✓
Custom plugin development	–	✓
Ongoing support + updates	✓	Priority support

Compare: a single ChatGPT Enterprise seat costs \$60/month. A 10-person team spends \$7,200/year and still sends data to the cloud. OpenClaw is a one-time investment with unlimited local usage forever.

What Makes This Different

You are not renting access to someone else's AI. You are buying a machine, pre-loaded with a battle-tested agent platform and plugins from the Organized AI Marketplace, configured specifically for your business. The hardware appreciates in utility as new models release — we push updates to your machine that make it more capable over time. And because everything runs locally, you get the one thing no cloud AI provider can offer: complete certainty that your business data stays inside your walls.

Common Questions

How fast is the test-drive setup?

Same day. Once your deposit clears, we provision your agent on our Mac Studio M3 Ultra and send you access credentials. Most clients are running their first tasks within hours.

Do I need technical expertise to use OpenClaw?

No. You interact with OpenClaw the same way you would text a colleague — via WhatsApp, a web interface, or API. The agent understands natural language and handles the technical complexity behind the scenes.

What happens to my data during the test drive?

Your data runs on our dedicated Mac Studio M3 Ultra during the test-drive period. It is not co-mingled with other clients. When your Mac Mini ships, we migrate everything to your hardware and wipe our copy.

What plugins come with my installation?

Your Mac Mini ships with curated plugins from the Organized AI Marketplace, selected based on your business vertical and usage during the test drive. The marketplace has 70+ vetted plugins across marketing, sales, product, data, dev, and martech — all tested for production use.

Can I add more capabilities after installation?

Yes. New plugins from the Organized AI Marketplace and updated models can be added remotely at any time. Our support plans include capability additions as part of the service.

What if a better model comes out?

We update your Mac Mini remotely. When a new open-source model outperforms your current stack, we download it, benchmark it against your workloads, and swap it in — all without disrupting your service.

Is this really unlimited usage?

Yes. There are no per-message fees, no token limits, and no rate throttling. The Mac Mini runs 24/7 and every inference is free after the hardware purchase. The only running cost is electricity, which averages around \$4–6 per month.

Can multiple team members use the same agent?

Yes. OpenClaw supports concurrent users. Your Mac Mini comfortably handles a team running different tasks simultaneously without performance degradation.

What if the machine has a hardware failure?

The Mac Mini is covered by Apple's standard warranty and optionally AppleCare+. Our support plan includes remote diagnostics. In the unlikely event of hardware failure, we can re-provision your agent on our Mac Studio M3 Ultra while the machine is repaired, so you experience zero downtime.

Ready to Get Started?

Your AI agent is ready to go. Place your deposit and start using OpenClaw today on our enterprise-grade infrastructure while we build your dedicated Mac Mini. No waiting. No setup hassle. No cloud dependency.

Next Step

Schedule a 30-minute walkthrough to see OpenClaw in action,
discuss your specific use case, and reserve your slot.

clawdbot.organizedai.vip

Organized AI

Building AI infrastructure that you own.